

What We've Learned About Our Customers — Walkthrough Script

Customer Segments & Cost to Serve, plus Autopay Removal Findings

Prepared by Momna Ali | For: Remy and Jonathan | Plain English throughout

Plain-English Summary

This covers two things. The main one, which Jonathan asked to be emphasized: some customers cost more to serve than others, and we can tell in advance which ones. Customers with lower credit scores call in 2.6 to 3 times more often than customers with higher credit scores, at every bill size. This is a real, measurable number for deciding which customers to bring on and how to plan for them.

The second part covers why customers turn off automatic bill payment (autopay). The team believed 38% of removals were new customers backing out early. The real number is 15.2%. Most removals (78%) happen quietly online with no phone call at all. About 2 out of 3 removals happen close to when a bill is due. The single biggest reason for phone-call removals is simply a customer's card expiring or being replaced (38%) — and the good news is, 90.5% of those customers turn autopay back on within 60 days, so it isn't a real problem, just a normal pause. Three other ideas were checked (payment timing, a surprise first-bill fee, and consent disputes) and all three are real, but rare.

Both pieces were built the same way: pulling real records directly from our own databases using SQL, then checking the numbers carefully before trusting them.

How to Walk Through It, Slide by Slide

Assume the audience has no background — explain everything as if for the first time.

Slide 1 — Title

Introduce both topics in one sentence — who costs more to serve, and why autopay gets turned off.

Say something like:

“Today I want to cover two things — what we've learned about which customers cost more to serve, and why customers turn off autopay.”

Slide 2 — Why We Looked At This (Part 1)

Start with the everyday truth behind this: every time a customer calls us, it costs money — an agent's time, wait times for other customers, and so on. Some customers almost never call. Others call again and again.

Explain the business question underneath: when we decide which customers to bring on, we mostly look at whether they'll pay their bills. But there's a second cost nobody was measuring — how much help they'll need after they join. Two customers paying the same bill can cost us very different amounts to take care of.

The goal: find out whether we can predict, ahead of time, which customers will need more help — using information we already have on day one, like their credit score.

Say something like:

“Every call costs us money, and some customers call far more than others. We wanted to know: can we tell on day one — before someone even becomes a customer — whether they're likely to be a high-support customer later? That matters for the cost-to-serve conversation, because right now we mostly judge customers on

whether they'll pay, not on how much help they'll need."

Slide 3 — How We Did This (Part 1)

Be specific about the data window if asked: this uses the last 6 months of real call records — every agent-handled call from our Texas residential customers.

Explain what went into it, simply: for every active customer, we pulled three things — their credit score (from our own customer records), their typical bill size, and a count of how many times they called us in those 6 months.

Explain the grouping: we sorted customers into a grid — credit score on one side (lower, meaning around 500 or below, versus higher, meaning 700 and up), bill size on the other (\$200–300, \$300–400, \$400+). Then, for each box in that grid, we calculated the average number of calls per customer.

One honest data-cleanup detail worth mentioning if asked: some customers show a credit score of zero in our system — that doesn't mean terrible credit, it means no score on file. We removed those from the comparison so they wouldn't distort the results.

The trust check: before believing any of this, we verified every group had thousands of customers in it — the smallest group still had 5,830 people. That means these are stable patterns, not flukes from a handful of unusual customers.

Say something like:

"We took the last 6 months of real call records, and for every customer we pulled three things: their credit score, their typical bill size, and how many times they called us. Then we sorted everyone into a grid — credit level on one side, bill size on the other — and calculated the average calls per customer in each box. And before trusting anything, we made sure every box had thousands of customers in it, so these aren't flukes."

Slide 4 — The Finding (Part 1) — THIS IS THE ONE TO EMPHASIZE

This is the slide Jonathan specifically wants emphasized — slow down, walk it carefully.

Walk the table row by row, out loud. Start with the smallest bills: "For customers with bills between \$200 and \$300, the lower-credit group averaged about 1 call each over 6 months. The higher-credit group averaged about 0.4 calls. That's 2.6 times more."

Then point out the pattern holds — and even grows — as bills get bigger: 2.7x at \$300–400, and 3x at \$400+.

Then make the key contrast explicit: you might expect bill size to be the thing that drives calls — bigger bill, more likely to call. But look across any single row: the bill size is the SAME in that row, and the difference is still huge. That's how we know it's the credit score doing the predicting, not the bill.

If she asks what 'lower' and 'higher' credit mean specifically: lower is roughly 500 and below, higher is roughly 700 and up — standard credit score ranges.

If she asks why lower-credit customers call more: we can see THAT they call more, clearly and consistently — but the data doesn't directly tell us WHY. Reasonable explanations include living closer to the financial edge (so any bill surprise hits harder), but that's interpretation, not something we've proven yet. Be honest about that line.

Land the business point: this gives a real, measurable number to a conversation that's mostly been instinct — which customers cost more to serve. It's relevant to onboarding standards, to deposit policies, and to deciding who gets proactive help.

Say something like:

"Let me walk through this table carefully, because it's the most important thing I'll show today. Take the first row — customers whose bills are \$200 to \$300. The lower-credit customers in that group called us about once each over six months. The higher-credit customers? Less than half a call each. Same bill size — 2.6 times more calls. And it gets stronger as bills grow: 3 times more at \$400-plus. The bill size isn't what's driving this — the credit score is. That's a real, measurable input into which customers cost more to serve."

Slide 5 — Why We Looked At This (Part 2)

Explain what autopay is first, simply: automatic bill payment — the customer's bill gets paid on its own each month without them doing anything. It's good for us (reliable payments) and good for them (no missed bills, no late fees).

State the belief being tested: the team believed about 38% of removals were brand-new customers backing out shortly after signing up — and if true, the fix would be about how we welcome new customers.

Why it matters to get this right: fixing the wrong problem wastes effort. If the real cause is something else entirely, we'd be improving onboarding while the actual issue continues untouched.

Say something like:

"Separately, we looked at autopay — automatic bill payment. The team believed about 38% of the customers turning it off were brand-new customers backing out early. If that were true, the fix would be about onboarding. We wanted to check whether it's actually true before anyone acts on it."

Slide 6 — How We Did This (Part 2)

Data window, if asked: the last 6 months of removals — every single one, about 33,000 in total.

Explain the channel split simply: our records show WHO made each change. If it was made through the website, the record shows the system did it — meaning the customer did it themselves online. If an agent's name is on it, it happened during a phone call.

Explain the call-reading part: for the phone-call removals, every call gets a short written summary of what happened. We searched those summaries for the actual words customers used — not categories we invented, but real phrases like 'card expired' or 'withdrawn earlier than expected.'

Explain the timing check: for every removal, we found that customer's next bill due date and measured how far away it was — to see whether removals cluster around bills coming due.

Say something like:

"We pulled every removal from the last 6 months — about 33,000 of them. Our records show who made each change, so we could split them into customers doing it themselves online versus during a phone call. For the phone calls, we read the written summaries of what customers actually said. And for all of them, we checked how close each removal was to that customer's next bill due date."

Slide 7 — The Original Belief Was Wrong

Point at the chart: assumed 38%, actual 15.2%.

If asked where the original 38% came from: it was an earlier rough estimate, not based on this kind of complete record-by-record check — this is the first time every removal was actually pulled and verified directly.

Say something like:

"The team believed 38% of removals were new customers. When we actually checked every record, the real number is 15.2% — much lower. Most removals come from long-time customers."

Slide 8 — How People Turn It Off

Point at the chart: 78% online, 22% by phone.

Draw out the implication: for nearly 4 out of 5 removals, we have no idea why — the customer just logged in and turned it off without telling anyone anything. Only the phone-call removals leave an explanation behind.

Say something like:

“Almost 4 out of 5 people just log in and turn it off themselves — no call, no explanation. That means for most removals, the only way to understand the ‘why’ is through the smaller group who did call.”

Slide 9 — When People Turn It Off

Point at the chart: both channels show about 2 in 3 removals within two weeks of the next bill due date.

Explain why this matters: even for the silent online removals where nobody tells us why, the timing gives us a clue — people act on autopay when a bill is about to land, not randomly through the month.

Say something like:

“Whether it’s a phone call or done online, about 2 out of 3 removals happen right around when a bill is due. Even for the silent online ones, that timing tells us something — people make this decision when a bill is about to hit.”

Slide 10 — The #1 Reason: A Changed Payment Card

The biggest, most surprising finding in the whole deck — take real time here.

Tell the story of how it was found: we didn’t start with this idea. We read a random sample of real call summaries with fresh eyes, and this pattern jumped out — call after call was simply ‘my card expired,’ ‘I lost my card,’ ‘I got a new card.’ Then we measured it across all the calls: 38%.

Then the crucial second step: we checked what happened NEXT for those customers. For every card-related removal, we looked at whether that same customer turned autopay back on within 60 days. Answer: 90.5% did.

Land the meaning clearly: this is not customers quitting autopay. It’s a routine pause while they sort out a new card — and it self-heals. The biggest ‘problem’ on the list turns out not to be a problem at all.

Say something like:

“This is the finding that surprised us most. We read a sample of real calls with fresh eyes, and one pattern kept appearing — ‘my card expired,’ ‘I got a new card.’ When we measured it properly, it was 38% of all phone-call removals — the single biggest reason, bigger than everything we’d guessed. And here’s the best part: 90.5% of those customers turn autopay back on within 60 days. So the biggest reason isn’t a complaint at all — it’s a routine card update that fixes itself.”

Slide 11 — Checking the Three Original Ideas

Give the backstory: these three came from real examples — Jonathan had shared two actual customer calls on Teams, and the third idea came from his broader complaint review. They weren’t guesses.

Walk each of the three, in plain terms, with its number:

- Payment date timing — the payment pulls on a day that doesn’t match when the customer gets paid (like the 1st, when their paycheck lands mid-month), sometimes causing other payments to bounce. Found in about 1 in 150 calls.

- Surprise first-bill fee — a brand-new customer gets an unexpected activation charge on a tiny first bill, making it look outrageous. About 1 in 1,000 calls.
- Never agreed to autopay — the customer says they never signed up for it in the first place. Also about 1 in 1,000 calls.

Close honestly: all three are real — we found genuine examples of each — but all three are small next to the card-change finding. The data isn't saying the examples were wrong; it's saying the biggest driver was something nobody had thought to check.

Say something like:

"We also checked the three ideas that came from the real examples Jonathan shared. All three are genuinely real — we found actual cases of each. The payment-timing issue shows up in about 1 in 150 calls, and the other two — the surprise first-bill fee and customers saying they never agreed — are each about 1 in 1,000. Real, but rare. The card-change finding is the much bigger story."

Slide 12 — Key Findings

Read all five rows as the full wrap-up of everything covered, both parts together.

Say something like:

"So putting it all together — credit score predicts who needs help, the old autopay belief was wrong, most people don't call us at all, the real #1 reason isn't a complaint, and the original three guesses are real but small."

Slide 13 — Next Steps

Frame as practical follow-ons across both topics.

Say something like:

"Here's what we'd suggest doing next with both of these."

Questions to Ask

1.

Does the credit-quality finding change how you're thinking about onboarding standards, or does it confirm what you already suspected?

2.

Should this same approach — grouping customers and counting real call patterns — be applied to other decisions, like which customers get extra support proactively?

3.

For the card-expiration finding — is a proactive reminder before expiration something worth prioritizing?

4.

Should we still loop in the Ops team on the consent issue, even though it turned out to be rare?

If They Ask... (Quick Answers)

Q: What time period does the customer-segment analysis cover?

A: The last 6 months of real call records — same window as the autopay work, so both analyses describe the same recent period.

Q: What exactly went into the analysis — what data was used?

A: Three things per customer, all from our own records: their credit score from our customer file, their typical bill size, and a count of their actual calls to us over 6 months. One cleanup detail: customers showing a credit score of zero were removed, because zero means 'no score on file,' not bad credit — including them would have distorted the comparison.

Q: How do you know the credit-score finding isn't just because lower-credit customers have bigger bills?

A: We specifically checked this by comparing customers at the SAME bill size, just different credit scores. The gap held at every bill size we checked — so it's really about credit score, not bill size.

Q: Is this fair to customers — are we treating people differently based on credit score?

A: This isn't about how customers are treated day to day. It's about understanding, in advance, which customers are more likely to need extra help, so we can plan for that — the same idea as reaching out before a bill surprises someone.

Q: How do you know the 90.5% re-enrollment number is right?

A: We checked, for every card-change-related removal, whether that same customer turned autopay back on within 60 days, using the real enrollment records directly. It's a direct count, not an estimate.

Q: Could these numbers change over time?

A: Yes, which is why we'd want to check them again periodically rather than treating this as a one-time answer.